



# Asynchronous Development

Lecture 5



# Asynchronous Development

- Concurrency
- Asynchronous Executor
- Future s
- Communication between tasks



# Concurrency

Preemptive and Cooperative



# Bibliography

for this section

**Brad Solomon**, *Async IO in Python: A Complete Walkthrough*



# Preemptive Concurrency

- MCUs are usually *single core*<sup>[1]</sup>
- Tasks in parallel require an OS<sup>[2]</sup>
- Tasks can be suspended at any time
- **Switching** the task is **expensive**
- Tasks that do a lot of I/O which makes the **switching time longer than the actual processing time**

1. RP2040 is a dual core MCU, we use only one core ↵
2. Running in an ISR is not considered a normal task ↵





# Cooperative Concurrency

- tasks cannot be interrupted<sup>[1]</sup>
- **hardware** works in an **asynchronous** way
- tasks cooperate
  - give up the MCU for other tasks to use it while they wait for hardware
- there is **no need for an OS**, everything is done in **one single flow**
- **no penalty** for saving and restoring the state

1. except for ISR ↔





# Asynchronous Executor

of Embassy



# Bibliography

for this section

**Embassy Documentation, *Embassy executor***





# Tasks

- `#[embassy_executor::main]`
  - starts the Embassy scheduler
  - defines the `main` task
- `#[embassy_executor::task]` - defines a new task
  - `pool_size` - is *optional* and defines how many identical tasks can be spawned
- the `main` task
  - initializes the `led`
  - spawns the `led_blink` task (adds to the scheduler)
  - uses `.await` to give up the MCU while waiting from the button

```
1  #[embassy_executor::task(pool_size = 2)]
2  async fn led_blink(mut led:Output<'static, PIN_X>) {
3      loop {
4          led.toggle();
5          Timer::after_secs(1).await;
6      }
7  }
8
9  #[embassy_executor::main]
10 async fn main(spawner: Spawner) {
11     // ...
12
13     // init led
14     spawner.spawn(led_blink(led)).unwrap();
15     info!("task started");
16
17     // init button
18     loop {
19         button.wait_for_rising_edge().await;
20         info!("button pressed");
21     }
22 }
```



# Tasks can stop the executor

- unless awaited, `async` functions are not executed
- tasks have to use `.await` in loops, otherwise they block the scheduler

```
1  #[embassy_executor::task]
2  async fn led_blink(mut led:Output<'static, PIN_X>) {
3      loop {
4          led.toggle();
5          // this does not execute anything
6          Timer::after_secs(1);
7          // infinite loop without `.await`
8          // that never gives up the MCU
9      }
10 }
11
12 #[embassy_executor::main]
13 async fn main(spawner: Spawner) {
14     // ..
15     loop {
16         button.wait_for_rising_edge().await;
17         info!("button pressed");
18     }
19 }
```



# How it works



- sleep when **all tasks wait for events**
- after an ISR is executed
  - if waiting for events, ask every task if it can execute (if the IRQ was what the task was `.await` ing for)
  - if a task is executing, continue the task until it `.await` s
- if a task never `.await` s, the executor does not run and never executes another task



# Priority Tasks



```
#[interrupt]
unsafe fn SWI_IRQ_1() {
    EXECUTOR_HIGH.on_interrupt()
}
#[interrupt]
unsafe fn SWI_IRQ_0() {
    EXECUTOR_MED.on_interrupt()
}
```

```
1 static EXECUTOR_HIGH: InterruptExecutor = InterruptExecutor::new();
2 static EXECUTOR_MED: InterruptExecutor = InterruptExecutor::new();
3 static EXECUTOR_LOW: StaticCell<Executor> = StaticCell::new();
4
5 #[entry]
6 fn main() -> ! {
7     // High-priority executor: SWI_IRQ_1, priority level 2
8     interrupt::SWI_IRQ_1.set_priority(Priority::P2);
9     let spawner = EXECUTOR_HIGH.start(interrupt::SWI_IRQ_1);
10    spawner.spawn(run_high()).unwrap();
11
12    // Medium-priority executor: SWI_IRQ_0, priority level 3
13    interrupt::SWI_IRQ_0.set_priority(Priority::P3);
14    let spawner = EXECUTOR_MED.start(interrupt::SWI_IRQ_0);
15    spawner.spawn(run_med()).unwrap();
16
17    // Low priority executor: runs in thread mode, using WFE/SEV
18    let executor = EXECUTOR_LOW.init(Executor::new());
19    executor.run(|spawner| {
20        unwrap!(spawner.spawn(run_low()));
21    });
22 }
```

priority executors run in ISRs, lower priority tasks are interrupted



# The Future type

a.k.a Promise in other languages



# Bibliography

for this section

**Bert Peters,** *How does async Rust work*



# Future

```
enum Poll<T> {  
    Pending,  
    Ready(T),  
}  
  
trait Future {  
    type Output;  
    fn poll(&mut self) -> Poll<Self::Output>;  
}
```

```
fn execute<F>(mut f: F) -> F::Output  
where  
    F: Future  
{  
    loop {  
        match f.poll() {  
            Poll::Pending => wait_for_event(),  
            Poll::Ready(value) => break value  
        }  
    }  
}
```





# Implementing a Future

```
1  enum SleepStatus {
2      SetAlarm,
3      WaitForAlarm,
4  }
5
6  struct Sleep {
7      timeout: usize,
8      status: SleepStatus,
9  }
10
11  impl Sleep {
12      pub fn new(timeout: usize) -> Sleep {
13          Sleep {
14              timeout,
15              status: SleepStatus::SetAlarm,
16          }
17      }
18  }
```

```
impl Future for Sleep {
    type Output = ();

    fn poll(&mut self) -> Poll<Self::Output> {
        match self.status {
            SleepStatus::SetAlarm => {
                ALARM.set_alarm(self.timeout);
                self.status = SleepStatus::WaitForAlarm;
                Poll::Pending
            }
            SleepStatus::WaitForAlarm => {
                if ALARM.expired() {
                    Poll::Ready(())
                } else {
                    Poll::Pending
                }
            }
        }
    }
}
```





# Executing Sleep

```
impl Future for Sleep {  
    type Output = ();  
  
    fn poll(&mut self) -> Poll<Self::Output> {  
        match self.status {  
            SleepStatus::SetAlarm => {  
                ALARM.set_alarm(self.timeout);  
                self.status = SleepStatus::WaitForAlarm;  
                Poll::Pending  
            }  
            SleepStatus::WaitForAlarm => {  
                if ALARM.expired() {  
                    Poll::Ready(())  
                } else {  
                    Poll::Pending  
                }  
            }  
        }  
    }  
}
```





# Async Rust

```
async fn blink(mut led: Output<'static, PIN_X>) {  
    led.on();  
    Timer::after_secs(1).await;  
    led.off();  
}
```

## Rust rewrites

```
struct Blink {  
    // status  
    status: BlinkStatus,  
    // local variables  
    led: Output<'static, PIN_X>,  
    timer: Option<impl Future>,  
}  
  
impl Blink {  
    pub fn new(led: Output<'static, PIN_X>) -> Blink {  
        Blink { status: BlinkStatus::Part1, led, timer: None }  
    }  
}  
  
fn blink(led: Output<'static, PIN_X>) -> Blink {  
    Blink::new(led)  
}
```

```
impl Future for Blink {  
    type Output = ();  
    fn poll(&mut self) -> Poll<Self::Output> {  
        loop {  
            match self.status {  
                BlinkStatus::Part1 => {  
                    self.led.on();  
                    self.timer1 = Some(Timer::after_secs(1));  
                    self.status = BlinkStatus::Part2;  
                }  
                BlinkStatus::Part2 => {  
                    if self.timer.unwrap().poll() == Poll::Pending {  
                        return Poll::Pending;  
                    } else {  
                        self.status = BlinkStatus::Part3;  
                    }  
                }  
                BlinkStatus::Part3 => {  
                    self.led.off();  
                    return Poll::Ready(());  
                }  
            }  
        }  
    }  
}
```



# Async Rust

- the Rust compiler rewrites `async` function into `Future`
- it does not know how to execute them
- executors are implemented into third party libraries

```
1  use engine::execute;
2
3  // Rust rewrites the function to a Future
4  async fn blink(mut led: Output<'static, PIN_X>) {
5      led.on();
6      Timer::after_secs(1).await;
7      led.off();
8  }
9
10 #[entry]
11 fn main() -> ! {
12     blink(); // this returns the Blink future, but does not execute it
13     blink().await; // does not work, as `main` is not an `async` function
14     execute(blink()); // this works, as `execute` executes the Blink future
15 }
```



# Executor

```
1  static TASKS: [Option<impl Future>; N] = [None, N];
2
3  fn executor() {
4      loop {
5          // ask all tasks to continue if they have available data
6          for task in TASKS.iter_mut() {
7              if let Some(task) = task {
8                  if Poll::Ready(_) = task.poll() {
9                      *task = None
10                 }
11             }
12         }
13
14         // wait for interrupts
15         cortex_m::asm::wfi();
16     }
17 }
```

- this is a simplified version, `Option<impl Future>` does not work
- the executor is not able to use `TASKS` like this
- an efficient executor will not poll all the tasks, it uses a `waker` that tasks use to signal the executor



# The Future trait

that Rust provides

```
1 trait Future {  
2     type Output;  
3  
4     fn poll(mut self: std::pin::Pin<&mut Self>, cx: &mut Context<'_>) -> Poll<Self::Output>;  
5 }
```

- `Pin` to `mut self`, which means that `self` cannot be moved
- `Context` which provides the `waker`
  - tasks are polled only if they ask the executor (by using the `wake` function)
- `embassy-rs` provides the execution engine





# Communication

between tasks



# Bibliography

for this section

**Omar Hiari**, *Sharing Data Among Tasks in Rust Embassy: Synchronization Primitives*



# Simultaneous Access

Rust forbids simultaneous writes access







# Exclusive Access

we want to sequentially access the resource





# Synchronization

safely share data between tasks

- `NoopMutex` - used for data shared between tasks within the **same executor**
- `CriticalSectionMutex` - used for data shared between multiple executors, ISRs and cores
- `ThreadModeMutex` - used for data shared between tasks within **low priority executors** (not running in **ISRs** mode) running on a **single core**



- ISRs are executed in parallel with tasks
- embassy allows registering priority executors, that run tasks in ISRs
- some MCUs have multiple cores



# Blocking Mutex

no `.await` allowed while the mutex is held

```
1  use embassy_sync::blocking_mutex::Mutex;
2
3  struct Data { /* ... */ }
4
5  static SHARED_DATA: Mutex<ThreadModeRawMutex, RefCell<Data>> = Mutex::new(RefCell::new(Data::new(/* ... */)));
6
7  #[embassy_executor::task]
8  async fn task1() {
9      // Load value from global context, modify and store
10     SHARED_DATA.lock(|f| {
11         let data = f.borrow_mut();
12         // edit data
13         f.replace(data);
14     });
15 }
```



# Async Mutex

`.await` is allowed while the Mutex is held, it will release the Mutex while `await` ing

```
1  use embassy_sync::mutex::Mutex;
2
3  struct Data { /* ... */ }
4
5  static SHARED: Mutex<ThreadModeRawMutex, Data> = Mutex::new(Data::new(/* ... */));
6
7  #[embassy_executor::task]
8  async fn task1() {
9      // Load value from global context, modify and store
10     {
11         let mut data = SHARED_DATA.lock().await;
12         // edit *data
13         Timer::after(Duration::from_millis(1000)).await;
14     }
15 }
```



# Channels

send data from a task to another

Embassy provides four types of channels synchronized using `Mutex` s

Type	Description
<code>Channel</code>	A Multiple Producer Multiple Consumer (MPMC) channel. Each message is only received by a single consumer.
<code>PriorityChannel</code>	A Multiple Producer Multiple Consumer (MPMC) channel. Each message is only received by a single consumer. Higher priority items are shifted to the front of the channel.
<code>Signal</code>	Signalling latest value to a single consumer.
<code>PubSubChannel</code>	A broadcast channel (publish-subscribe) channel. Each message is received by all consumers.



# Channel and Signal

sends data from one task to another

**Channel** - A Multiple Producer Multiple Consumer (MPMC) channel. Each message is only received by a single consumer.

**Signal** - Signalling latest value to a single consumer.





# PriorityChannel

sends data from one task to another with a priority

`PriorityChannel` - A Multiple Producer Multiple Consumer (MPMC) channel. Each message is only received by a single consumer. Higher priority items are shifted to the front of the channel.





# PubSubChannel

sends data from one task to all receiver tasks

`PubSubChannel` - A broadcast channel (publish-subscribe) channel. Each message is received by all consumers.







# Channel Example

```
1  enum LedState { On, Off }
2  static CHANNEL: Channel<ThreadModeRawMutex, LedState, 64> = Channel::new();
3
4  #[embassy_executor::main]
5  async fn main(spawner: Spawner) {
6      // init led
7      spawner.spawn(execute_led(CHANNEL.sender(), Duration::from_millis(500))));
8      loop {
9          match CHANNEL.receive().await {
10             LedState::On => led.on(),
11             LedState::Off => led.off()
12          }
13      }
14  }
15
16  #[embassy_executor::task]
17  async fn execute_led(control: Sender<'static, ThreadModeRawMutex, LedState, 64>, delay: Duration) {
18      let mut ticker = Ticker::every(delay);
19      loop {
20          control.send(LedState::On).await;
21          ticker.next().await;
22          control.send(LedState::Off).await;
23          ticker.next().await;
24      }
25  }
```



# Conclusion

we discussed about

- Preemptive & Cooperative Concurrency
- Asynchronous Executor
- `Future` s and how Rust rewrites `async` function
- Communication between tasks